

Results

August 30, 2020

1 MARS-Gym Experiments - Trivago Results

```
[1]: import sys, os
os.environ['PATH_EVALUATION'] = "../..output/evaluation/"
os.environ['PATH_EVAL_REINFORCEMENT'] = "../..output/interaction/"
os.environ['PATH_TRAIN'] = "../..output/models/"

#import plotly.graph_objects as go
import plotly.graph_objects as go
from plotly.offline import init_notebook_mode, iplot
init_notebook_mode(connected=True)
import numpy as np
import pandas as pd
import os
import json
from pandas.io.json import json_normalize
import matplotlib.pyplot as plt
import seaborn as sns

from mars_gym.tools.eval_viz.app import fetch_interaction_results_path,
↳load_all_interaction_metrics, fetch_results_path, load_data_metrics,
↳load_eval_params, filter_df, load_fairness_metrics, load_fairness_df
from mars_gym.tools.eval_viz.plot import plot_line_interaction,
↳plot_exploration_arm, get_colors, confidence, plot_bar
from mars_gym.tools.eval_viz.app import PATH_EVAL_REINFORCEMENT

from util import list_paths_per_model, load_interactions_params,
↳get_metrics_reward, group_metrics, load_dataset, plot_cum_reward,
↳plot_fairness_metrics, plot_fairness_mistreatment, plot_fairness_treatment
```

/media/workspace/DeepFood/deep-reco-gym/src/mars_gym/tools/eval_viz/app.py:13:
FutureWarning:

Passing a negative integer is deprecated in version 1.0 and will not be supported in future version. Instead, use None to not limit the column width.

[]:

This notebook includes all viz of paper ‘MARS-Gym: A Gym framework to model, train, and evaluate recommendation systems for marketplaces’ and can be used for reproducibility or example of framework usage.

We need run some scripts before vizualise results: `* scripts/simulation/chicago_usa_script.sh * scripts/metrics/metrics_chicago_usa_script.sh * scripts/metrics/fairness_recsys_script.sh`

1.1 Load Results

```
[2]: result_paths = list_paths_per_model(PATH_EVAL_REINFORCEMENT)
result_paths[:5]
```

```
[2]: ['../output/interaction/InteractionTraining/results/InteractionTraining___ma
rs_gym_model_b___alpha___1e_05__a06cfb73a9',
'../output/interaction/InteractionTraining/results/InteractionTraining___ma
rs_gym_model_b___alpha___1e_05__f9e2903219',
'../output/interaction/InteractionTraining/results/InteractionTraining___ma
rs_gym_model_b___exploration_th_6822e93237',
'../output/interaction/InteractionTraining/results/InteractionTraining___ma
rs_gym_model_b___explore_rounds_0b659e17e7',
'../output/interaction/InteractionTraining/results/InteractionTraining___ma
rs_gym_model_b___epsilon___0_1__76d6a2cc36']
```

```
[3]: # We will only filter the city of Chicago to exemplify graphs and results.
#

#city = "Rio de Janeiro, Brazil"
#city = "Como, Italy"
city = "Chicago, USA"
#city = "New York, USA"
#city = 'Copenhagen, Denmark'
#city='recsys'

df_params = load_interactions_params(result_paths)
df_params = df_params[df_params['data_frames_preparation_extra_params.
↪filter_city'] == city]
df_params.head()
```

/media/workspace/DeepFood/mars-gym-experiments/scripts/notebooks/util.py:62:
FutureWarning:

pandas.io.json.json_normalize is deprecated, use pandas.json_normalize instead

```
[3]:          project  minimum_interactions  session_test_size  \
0  trivago.config.trivago_experiment      5                0.1
```

```

0   trivago.config.trivago_experiment  5           0.1
0   trivago.config.trivago_experiment  5           0.1
0   trivago.config.trivago_experiment  5           0.1
0   trivago.config.trivago_experiment  5           0.1

    dataset_split_method  val_size  n_splits  split_index  sampling_strategy \
0   time                  0.2        5         0         none
0   time                  0.2        5         0         none
0   time                  0.2        5         0         none
0   time                  0.2        5         0         none
0   time                  0.2        5         0         none

    balance_fields  use_sampling_in_validation  ...  bandit_policy_params.alpha \
0   []             False                      ...  0.00001
0   []             False                      ...  0.00001
0   []             False                      ...  NaN
0   []             False                      ...  NaN
0   []             False                      ...  NaN

                                interaction \
0   ../../output/interaction/InteractionTraining/results/InteractionTraining___m
ars_gym_model_b___alpha___1e_05__a06cfb73a9
0   ../../output/interaction/InteractionTraining/results/InteractionTraining___m
ars_gym_model_b___alpha___1e_05__f9e2903219
0   ../../output/interaction/InteractionTraining/results/InteractionTraining___m
ars_gym_model_b___exploration_th_6822e93237
0   ../../output/interaction/InteractionTraining/results/InteractionTraining___m
ars_gym_model_b___explore_rounds_0b659e17e7
0   ../../output/interaction/InteractionTraining/results/InteractionTraining___m
ars_gym_model_b___epsilon___0_1__76d6a2cc36

    bandit_policy_params.exploration_threshold  bandit_policy_params.decay_rate \
0   NaN                                         NaN
0   NaN                                         NaN
0   0.7                                         0.000097
0   NaN                                         0.000187
0   NaN                                         NaN

    bandit_policy_params.explore_rounds  bandit_policy_params.epsilon \
0   NaN                                  NaN
0   NaN                                  NaN
0   NaN                                  NaN
0   1000.0                              NaN
0   NaN                                  0.1

    loss_function_params.clip  bandit_policy_params.logit_multiplier \
0   NaN                       NaN

```

	bandit_policy_params.arg	bandit_policy_params.v_sq
0 NaN		NaN
0 NaN		NaN
0 NaN		NaN
0 NaN		NaN
0 NaN		NaN

```
[4]: # Group results
```

```
[4]: bandit_policy_class observation \
```

```
0 [../output/interaction/InteractionTraining/results/InteractionTraining____mars_gym_model_b__exploration_th_6822e93237, ../output/interaction/InteractionTraining/results/InteractionTraining____mars_gym_model_b__exploration_th_83ef3f55d3, ../output/interaction/InteractionTraining/results/InteractionTraining____mars_gym_model_b__exploration_th_42716ae133, ../output/interaction/InteractionTraining/results/InteractionTraining____mars_gym_model_b__exploration_th_d3663f58fc, ../output/interaction/InteractionTraining/results/InteractionTraining____mars_gym_model_b__exploration_th_520cebc570, ../output/interaction/InteractionTraining/results/InteractionTraining____mars_gym_model_b__exploration_th_d70af8a008, ../output/interaction/InteractionTraining/results/InteractionTraining____mars_gym_model_b__exploration_th_6b44643887, ../output/interaction/InteractionTraining/results/InteractionTraining____mars_gym_model_b__exploration_th_2cc29714f4]
```

4

```

onTraining/results/InteractionTraining___mars_gym_model_b___alpha___1e_05__75ec
cc229a, ../../output/interaction/InteractionTraining/results/InteractionTraining
___mars_gym_model_b___alpha___1e_05__bad81a6465, ../../output/interaction/Inter
actionTraining/results/InteractionTraining___mars_gym_model_b___alpha___1e_05__
debd700fb8, ../../output/interaction/InteractionTraining/results/InteractionTrai
ning___mars_gym_model_b___alpha___1e_05__2165771e8e, ../../output/interaction/I
nteractionTraining/results/InteractionTraining___mars_gym_model_b___alpha___1e_
05__7966e4e52b, ../../output/interaction/InteractionTraining/results/Interaction
Training___mars_gym_model_b___alpha___1e_05__0b623b96b8, ../../output/interacti
on/InteractionTraining/results/InteractionTraining___mars_gym_model_b___alpha__
_1e_05__d5981787b1]

```

name \

```

0 mars_gym.model.bandit.AdaptiveGreedy
1 mars_gym.model.bandit.CustomRewardModelLinUCB

```

model \

```

0 ../../output/interaction/InteractionTraining/results/InteractionTraining___m
ars_gym_model_b___exploration_th_d3663f58fc
1 ../../output/interaction/InteractionTraining/results/InteractionTraining___m
ars_gym_model_b___alpha___1e_05__7966e4e52b

```

model_list

```

0 [../../output/interaction/InteractionTraining/results/InteractionTraining___
mars_gym_model_b___exploration_th_6822e93237, ../../output/interaction/Interacti
onTraining/results/InteractionTraining___mars_gym_model_b___exploration_th_83ef
3f55d3, ../../output/interaction/InteractionTraining/results/InteractionTraining
___mars_gym_model_b___exploration_th_42716ae133, ../../output/interaction/Inter
actionTraining/results/InteractionTraining___mars_gym_model_b___exploration_th_
d3663f58fc, ../../output/interaction/InteractionTraining/results/InteractionTrai
ning___mars_gym_model_b___exploration_th_520cebc570, ../../output/interaction/I
nteractionTraining/results/InteractionTraining___mars_gym_model_b___exploration
_th_d70af8a008, ../../output/interaction/InteractionTraining/results/Interaction
Training___mars_gym_model_b___exploration_th_6b44643887, ../../output/interacti
on/InteractionTraining/results/InteractionTraining___mars_gym_model_b___explora
tion_th_2cc29714f4]
1 [../../output/interaction/InteractionTraining/results/InteractionTraining___
mars_gym_model_b___alpha___1e_05__ef475653f1, ../../output/interaction/Interacti
onTraining/results/InteractionTraining___mars_gym_model_b___alpha___1e_05__75ec
cc229a, ../../output/interaction/InteractionTraining/results/InteractionTraining
___mars_gym_model_b___alpha___1e_05__bad81a6465, ../../output/interaction/Inter
actionTraining/results/InteractionTraining___mars_gym_model_b___alpha___1e_05__
debd700fb8, ../../output/interaction/InteractionTraining/results/InteractionTrai
ning___mars_gym_model_b___alpha___1e_05__2165771e8e, ../../output/interaction/I
nteractionTraining/results/InteractionTraining___mars_gym_model_b___alpha___1e_
05__7966e4e52b, ../../output/interaction/InteractionTraining/results/Interaction
Training___mars_gym_model_b___alpha___1e_05__0b623b96b8, ../../output/interacti

```

```
on/InteractionTraining/results/InteractionTraining___mars_gym_model_b___alpha__
_1e_05__d5981787b1]
```

```
[ ]:
```

```
[5]: reward_metrics = []
for i, row in list_bandits.iterrows():
    df_metrics = get_metrics_reward(row['interaction'])
    df_metrics['bandit'] = row['name']
    df_metrics['model'] = row['model']

    df_g_metrics = group_metrics(df_metrics)
    reward_metrics.append(df_g_metrics)
```

```
[6]: df_all = pd.concat(reward_metrics)
df_all = df_all.reset_index().sort_values('bandit', ascending=False).
    ↪set_index('bandit')
df_all
```

```
[6]:
```

	r_mean mean	r_mean std \
bandit		
mars_gym.model.bandit.SoftmaxExplorer	0.275622	0.013914
mars_gym.model.bandit.RandomPolicy	0.040749	NaN
mars_gym.model.bandit.PercentileAdaptiveGreedy	0.284485	0.010725
mars_gym.model.bandit.LinUCB	0.121703	0.007030
mars_gym.model.bandit.LinThompsonSampling	0.048501	NaN
mars_gym.model.bandit.FixedPolicyPopularItem	0.104567	0.000000
mars_gym.model.bandit.ExploreThenExploit	0.228906	0.006917
mars_gym.model.bandit.EpsilonGreedy	0.285147	0.015851
mars_gym.model.bandit.CustomRewardModelLinUCB	0.297863	0.024824
mars_gym.model.bandit.AdaptiveGreedy	0.215609	0.030293

	r_reward mean	r_reward std \
bandit		
mars_gym.model.bandit.SoftmaxExplorer	2915.250000	147.173124
mars_gym.model.bandit.RandomPolicy	431.000000	NaN
mars_gym.model.bandit.PercentileAdaptiveGreedy	3009.000000	113.439726
mars_gym.model.bandit.LinUCB	1287.250000	74.351962
mars_gym.model.bandit.LinThompsonSampling	513.000000	NaN
mars_gym.model.bandit.FixedPolicyPopularItem	1106.000000	0.000000
mars_gym.model.bandit.ExploreThenExploit	2421.142857	73.156063
mars_gym.model.bandit.EpsilonGreedy	3016.000000	167.654406
mars_gym.model.bandit.CustomRewardModelLinUCB	3150.500000	262.564822
mars_gym.model.bandit.AdaptiveGreedy	2280.500000	320.404209

	r_reward count \
bandit	

<code>mars_gym.model.bandit.SoftmaxExplorer</code>	8
<code>mars_gym.model.bandit.RandomPolicy</code>	1
<code>mars_gym.model.bandit.PercentileAdaptiveGreedy</code>	8
<code>mars_gym.model.bandit.LinUCB</code>	8
<code>mars_gym.model.bandit.LinThompsonSampling</code>	1
<code>mars_gym.model.bandit.FixedPolicyPopularItem</code>	2
<code>mars_gym.model.bandit.ExploreThenExploit</code>	7
<code>mars_gym.model.bandit.EpsilonGreedy</code>	8
<code>mars_gym.model.bandit.CustomRewardModelLinUCB</code>	8
<code>mars_gym.model.bandit.AdaptiveGreedy</code>	8

`model|first`

`bandit`

<code>mars_gym.model.bandit.SoftmaxExplorer</code>	<code>../../output/interaction/InteractionTraining/results/InteractionTraining___mars_gym_model_b___logit_multipli_34358aae37</code>
<code>mars_gym.model.bandit.RandomPolicy</code>	<code>../../output/interaction/InteractionTraining/results/InteractionTraining___mars_gym_model_b___1649c7bd6b</code>
<code>mars_gym.model.bandit.PercentileAdaptiveGreedy</code>	<code>../../output/interaction/InteractionTraining/results/InteractionTraining___mars_gym_model_b___exploration_th_36a2721c50</code>
<code>mars_gym.model.bandit.LinUCB</code>	<code>../../output/interaction/InteractionTraining/results/InteractionTraining___mars_gym_model_b___alpha___1e_05__335eb27eb5</code>
<code>mars_gym.model.bandit.LinThompsonSampling</code>	<code>../../output/interaction/InteractionTraining/results/InteractionTraining___mars_gym_model_b___v_sq___0_1__4cd65ca01d</code>
<code>mars_gym.model.bandit.FixedPolicyPopularItem</code>	<code>../../output/interaction/InteractionTraining/results/InteractionTraining___mars_gym_model_b___arg___2__9cb6645207</code>
<code>mars_gym.model.bandit.ExploreThenExploit</code>	<code>../../output/interaction/InteractionTraining/results/InteractionTraining___mars_gym_model_b___explore_rounds_a180ca8b8f</code>
<code>mars_gym.model.bandit.EpsilonGreedy</code>	<code>../../output/interaction/InteractionTraining/results/InteractionTraining___mars_gym_model_b___epsilon___0_1__4645466ef3</code>
<code>mars_gym.model.bandit.CustomRewardModelLinUCB</code>	<code>../../output/interaction/InteractionTraining/results/InteractionTraining___mars_gym_model_b___alpha___1e_05__7966e4e52b</code>
<code>mars_gym.model.bandit.AdaptiveGreedy</code>	<code>../../output/interaction/InteractionTraining/results/InteractionTraining___mars_gym_model_b___exploration_th_d3663f58fc</code>

1.2 DataViz - Plot Images

```
[7]: #list_bandits['name'] = list_bandits['name'].apply(lambda x: "popular_item" if
    ↪ x == "PopularItem" else x)
df_models = list_bandits.set_index('name')[['model_list']]
df_models.index
```

```
[7]: Index(['mars_gym.model.bandit.AdaptiveGreedy',
    'mars_gym.model.bandit.CustomRewardModelLinUCB',
    'mars_gym.model.bandit.EpsilonGreedy',
    'mars_gym.model.bandit.ExploreThenExploit',
    'mars_gym.model.bandit.FixedPolicyPopularItem',
    'mars_gym.model.bandit.LinThompsonSampling',
    'mars_gym.model.bandit.LinUCB',
    'mars_gym.model.bandit.PercentileAdaptiveGreedy',
    'mars_gym.model.bandit.RandomPolicy',
    'mars_gym.model.bandit.SoftmaxExplorer'],
    dtype='object', name='name')
```

```
[8]: sample_size      = 15000
window_size          = 1000
bandits              = df_models.index

df_group = load_dataset(df_models, bandits, sample_size)
df_group.head()
```

```
[8]:
```

						interaction	idx \
0	InteractionTraining___mars_gym_model_b___exploration_th_2cc29714f4						0
1	InteractionTraining___mars_gym_model_b___exploration_th_2cc29714f4						1
2	InteractionTraining___mars_gym_model_b___exploration_th_2cc29714f4						2
3	InteractionTraining___mars_gym_model_b___exploration_th_2cc29714f4						3
4	InteractionTraining___mars_gym_model_b___exploration_th_2cc29714f4						4

	user	item	reward	ps	index_env	_idx	mean_reward \
0	559	191	0.0	0.975610	10	1	0.0
1	1118	486	0.0	1.951220	43	2	0.0
2	1677	644	0.0	2.926829	96	3	0.0
3	1908	1263	0.0	3.902439	186	4	0.0
4	1971	1397	0.0	4.878049	297	5	0.0

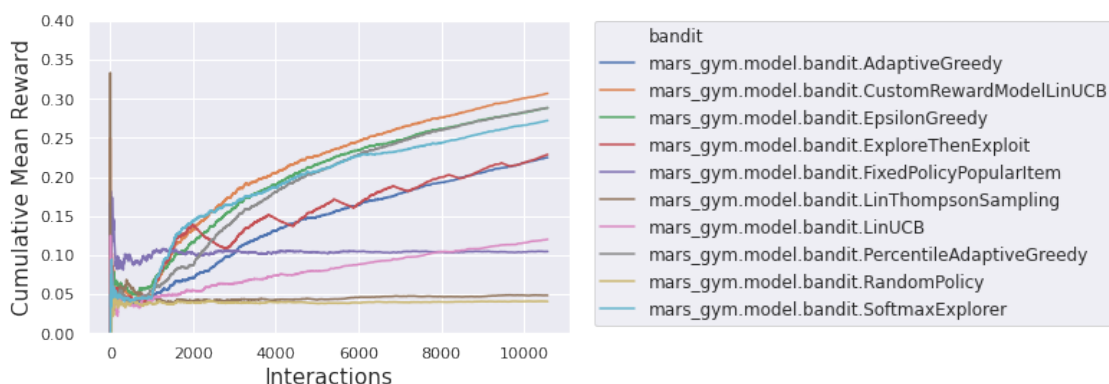
	bandit
0	mars_gym.model.bandit.AdaptiveGreedy
1	mars_gym.model.bandit.AdaptiveGreedy
2	mars_gym.model.bandit.AdaptiveGreedy
3	mars_gym.model.bandit.AdaptiveGreedy
4	mars_gym.model.bandit.AdaptiveGreedy

1.2.1 Bandit Simulation Results.

We performed the simulation of each proposed task until the convergence of most methods to observe cumulative mean reward throughout the simulation. Bandit Simulation Results for “Chicago, USA” task.

```
[9]: df_group_mean = df_group.groupby(['bandit', 'idx']).median().reset_index()

plot_cum_reward(df_group_mean, legend='full', ylim=0.4)
```



```
[ ]:
```

1.2.2 Recommendation Metrics and Off-Policy Evaluation

We evaluated the bandits according to traditional recommendation metrics and off-policy metrics in the test subset of the “Chicago, USA” task

```
[10]: input_metrics = [
    "precision_at_1",
    "ndcg_at_5",
    "coverage_at_5",
    "personalization_at_5",
    "IPS",
    "SNIPS",
    "DirectEstimator",
    "DoublyRobust"
]
```

```
[11]: # Load Métrics and Params
#
input_models_eval = fetch_results_path().keys()
data_metrics      = load_data_metrics()
data_params       = load_eval_params()
input_params      = []
```

```
df_metrics      = filter_df(data_metrics, input_models_eval, input_metrics,
    ↪ "")
df_eval_params  = filter_df(data_params, input_models_eval, input_params).
    ↪ transpose()
```

/media/workspace/DeepFood/deep-reco-gym/src/mars_gym/tools/eval_viz/util.py:44:
FutureWarning:

pandas.io.json.json_normalize is deprecated, use pandas.json_normalize instead

```
[12]: # Calculate confidence metrics
#
confidence_metrics = data_metrics[[c for c in data_metrics.columns if "_C" in
    ↪ c]]
for c in input_metrics:
    c_column = c + "_C"
    confidence_metrics[c_column] = data_metrics[c_column] if c_column in
    ↪ data_metrics else None
confidence_metrics = confidence_metrics[[c + "_C" for c in input_metrics]]
```

```
[13]: # Merge and Mean
#
df_metrics['model_task'] = data_metrics.model_task
df_params['model_task'] = df_params.interaction.apply(lambda x: x.split("/")
    ↪)[-1])

df_result = df_metrics.merge(
    df_params[['bandit_policy_class', 'model_task']].reset_index()
    , on='model_task')

df_result = df_result.groupby(['bandit_policy_class']).mean()
df_result
```

```
[13]:
```

bandit_policy_class	precision_at_1	ndcg_at_5	\
mars_gym.model.bandit.AdaptiveGreedy	0.317639	0.403579	
mars_gym.model.bandit.CustomRewardModelLinUCB	0.328249	0.442917	
mars_gym.model.bandit.EpsilonGreedy	0.302056	0.443252	
mars_gym.model.bandit.ExploreThenExploit	0.307692	0.418607	
mars_gym.model.bandit.FixedPolicy	0.074271	0.171294	
mars_gym.model.bandit.LinThompsonSampling	0.039788	0.137056	
mars_gym.model.bandit.LinUCB	0.075928	0.207314	
mars_gym.model.bandit.PercentileAdaptiveGreedy	0.336870	0.438648	
mars_gym.model.bandit.RandomPolicy	0.039788	0.138404	
mars_gym.model.bandit.SoftmaxExplorer	0.301724	0.452869	

	coverage_at_5 \
bandit_policy_class	
mars_gym.model.bandit.AdaptiveGreedy	0.390850
mars_gym.model.bandit.CustomRewardModelLinUCB	0.362862
mars_gym.model.bandit.EpsilonGreedy	0.343287
mars_gym.model.bandit.ExploreThenExploit	0.333071
mars_gym.model.bandit.FixedPolicy	0.373850
mars_gym.model.bandit.LinThompsonSampling	0.423900
mars_gym.model.bandit.LinUCB	0.271450
mars_gym.model.bandit.PercentileAdaptiveGreedy	0.375563
mars_gym.model.bandit.RandomPolicy	0.389600
mars_gym.model.bandit.SoftmaxExplorer	0.330988

	personalization_at_5 \
bandit_policy_class	
mars_gym.model.bandit.AdaptiveGreedy	0.767565
mars_gym.model.bandit.CustomRewardModelLinUCB	0.729376
mars_gym.model.bandit.EpsilonGreedy	0.733547
mars_gym.model.bandit.ExploreThenExploit	0.731641
mars_gym.model.bandit.FixedPolicy	0.759891
mars_gym.model.bandit.LinThompsonSampling	0.771211
mars_gym.model.bandit.LinUCB	0.696228
mars_gym.model.bandit.PercentileAdaptiveGreedy	0.743905
mars_gym.model.bandit.RandomPolicy	0.775653
mars_gym.model.bandit.SoftmaxExplorer	0.725512

	IPS	SNIPS \
bandit_policy_class		
mars_gym.model.bandit.AdaptiveGreedy	0.299046	0.308466
mars_gym.model.bandit.CustomRewardModelLinUCB	0.305576	0.315562
mars_gym.model.bandit.EpsilonGreedy	0.296847	0.295397
mars_gym.model.bandit.ExploreThenExploit	0.296997	0.294139
mars_gym.model.bandit.FixedPolicy	0.076116	0.077219
mars_gym.model.bandit.LinThompsonSampling	0.037423	0.035373
mars_gym.model.bandit.LinUCB	0.053439	0.056094
mars_gym.model.bandit.PercentileAdaptiveGreedy	0.321505	0.317082
mars_gym.model.bandit.RandomPolicy	0.041457	0.040577
mars_gym.model.bandit.SoftmaxExplorer	0.287223	0.287718

	DirectEstimator	DoublyRobust \
bandit_policy_class		
mars_gym.model.bandit.AdaptiveGreedy	0.200657	0.267138
mars_gym.model.bandit.CustomRewardModelLinUCB	0.199502	0.266484
mars_gym.model.bandit.EpsilonGreedy	0.187270	0.254888
mars_gym.model.bandit.ExploreThenExploit	0.190652	0.256225
mars_gym.model.bandit.FixedPolicy	0.084967	0.077790

<code>mars_gym.model.bandit.LinThompsonSampling</code>	0.041869	0.039049
<code>mars_gym.model.bandit.LinUCB</code>	0.054993	0.051205
<code>mars_gym.model.bandit.PercentileAdaptiveGreedy</code>	0.198404	0.280774
<code>mars_gym.model.bandit.RandomPolicy</code>	0.042541	0.041919
<code>mars_gym.model.bandit.SoftmaxExplorer</code>	0.188744	0.253449

	index
<code>bandit_policy_class</code>	
<code>mars_gym.model.bandit.AdaptiveGreedy</code>	0
<code>mars_gym.model.bandit.CustomRewardModelLinUCB</code>	0
<code>mars_gym.model.bandit.EpsilonGreedy</code>	0
<code>mars_gym.model.bandit.ExploreThenExploit</code>	0
<code>mars_gym.model.bandit.FixedPolicy</code>	0
<code>mars_gym.model.bandit.LinThompsonSampling</code>	0
<code>mars_gym.model.bandit.LinUCB</code>	0
<code>mars_gym.model.bandit.PercentileAdaptiveGreedy</code>	0
<code>mars_gym.model.bandit.RandomPolicy</code>	0
<code>mars_gym.model.bandit.SoftmaxExplorer</code>	0

```
[14]: df_result.round(3).to_csv('metrics.csv')
```

```
[15]: # fig = plot_bar(df_result, confidence = None, title="Comparison of Recsys
↳ Metrics")

# fig.update_layout(xaxis_showgrid=False, yaxis_showgrid=False)
# fig.update_layout(font={'family': 'Courier New, monospace', 'size': 14},
↳ height=550)
# fig#.show(renderer="svg")

# #fig.write_image("output/tmp/metrics/recsys_rank.eps")
```

```
[ ]:
```

1.2.3 Fairness Results

We evaluated the SoftmaxExplorer bandit on the “RecSys Cities” task, in the perspective of disparate mistreatment and disparate treatment. We selected a few attributes that we judged to be sensitive for all partners in the marketplace.

```
[16]: # task_id:
↳ InteractionTraining___mars_gym_model_b___logit_multipli_9dd8714dfd_96b9d686ee
```

```
[ ]:
```

Feature - Accessibility

```
[17]: input_models_eval = _
      ↪ 'InteractionTraining___mars_gym_model_b___logit_multipli_9dd8714dfd_96b9d686ee'
      list_input_features = ['accessible parking', 'accessible hotel']
      input_metrics = 'true_positive_rate'
```

```
[18]: # fig, df_metrics, df_total = plot_fairness_metrics(input_models_eval, _
      ↪ input_features, input_metrics, sub_map = None)
      metrics = []
      sub_map = None

      for input_features in list_input_features:
          df_all_metrics = load_fairness_metrics().loc[input_models_eval]
          df_instances = load_fairness_df().loc[input_models_eval]
          df_all_metric_filter = df_all_metrics[df_all_metrics.sub_key.
          ↪ isin([input_features])]

          columns = list(np.unique(['sub_key', 'sub', 'feature', _
          ↪ 'total_class', 'total_individuals'] + [input_metrics]))
          if input_metrics+"_C" in df_all_metrics.columns:
              columns.append(input_metrics+"_C")

          df_metrics = filter_df(df_all_metrics, input_models_eval, columns, _
          ↪ 'sub')

          df_metrics = df_metrics[df_metrics.sub_key.isin([input_features])]
          if sub_map != None:
              df_metrics['sub'] = df_metrics['sub'].map(sub_map)

          df_metrics = df_metrics.sort_values("feature").set_index("sub")
          df_metrics = df_metrics.sort_values(input_metrics)
          metrics.append(df_metrics)

      metrics = pd.concat(metrics)
      metrics
```

```
[18]:
```

	feature	sub_key	total_class	total_individuals \
sub				
1	accessible parking.1	accessible parking	1063	2380
0	accessible parking.0	accessible parking	3193	10038
1	accessible hotel.1	accessible hotel	1345	3799
0	accessible hotel.0	accessible hotel	3114	8619
	true_positive_rate	true_positive_rate_C		
sub				
1	0.222664	0.023616		
0	0.297201	0.013806		
1	0.242414	0.020492		

```
0    0.287845    0.014361
```

```
[19]: metrics.reset_index()
metrics['sub'] = ['parking<br>True', 'parking<br>False', 'hotel<br>True',
↳ 'hotel<br>False']
metrics['color'] = [2,2, 5,5]
metrics = metrics.set_index('sub')
metrics
```

```
[19]:
```

	feature	sub_key	total_class	\
sub				
parking True	accessible parking.1	accessible parking	1063	
parking False	accessible parking.0	accessible parking	3193	
hotel True	accessible hotel.1	accessible hotel	1345	
hotel False	accessible hotel.0	accessible hotel	3114	

	total_individuals	true_positive_rate	true_positive_rate_C	\
sub				
parking True	2380	0.222664	0.023616	
parking False	10038	0.297201	0.013806	
hotel True	3799	0.242414	0.020492	
hotel False	8619	0.287845	0.014361	

	color
sub	
parking True	2
parking False	2
hotel True	5
hotel False	5

```
[20]: fig = plot_fairness_mistreatment(metrics, input_metrics, title="")

fig.update_layout(shapes=[
    dict(
        type= 'line',
        line=dict(
            width=1,
            dash="dot",
        ),
        xref='paper', x0= 0, x1= 1,
        yref='y',      y0= 1.5, y1= 1.5
    )
])

# Set custom x-axis labels
fig.update_xaxes(nticks=4)
```

```

fig.update_layout(yaxis={'categoryorder':'array', 'categoryarray':
    ↳['parking<br>False', 'parking<br>True', 'hotel<br>False', 'hotel<br>True']})
fig.update_layout(xaxis_showgrid=False, yaxis_showgrid=False,
    ↳yaxis_title="Accessible")
fig.update_layout(font={'family': 'Courier New, monospace', 'size': 22},
    ↳height=550)

fig.show()
#fig.write_image("output/tmp/metrics/paper/recsys_fairnes_{}.eps".
    ↳format("accessible"))
#fig.show(renderer="svg")

```

[]:

Feature - Accommodation Type

```

[22]: input_models_eval =
    ↳'InteractionTraining___mars_gym_model_b___logit_multipli_9dd8714dfd_96b9d686ee'
list_input_features    = ['hotel', 'house / apartment']
input_metrics          = 'true_positive_rate'

```

```

[23]: metrics = []
for input_features in list_input_features:
    df_all_metrics    = load_fairness_metrics().loc[input_models_eval]
    df_instances      = load_fairness_df().loc[input_models_eval]
    df_all_metric_filter = df_all_metrics[df_all_metrics.sub_key.
    ↳isin([input_features])]

    columns           = list(np.unique(['sub_key', 'sub', 'feature',
    ↳'total_class', 'total_individuals'] + [input_metrics]))
    if input_metrics+"_C" in df_all_metrics.columns:
        columns.append(input_metrics+"_C")

    df_metrics        = filter_df(df_all_metrics, input_models_eval, columns,
    ↳'sub')

    df_metrics        = df_metrics[df_metrics.sub_key.isin([input_features])]
    if sub_map != None:
        df_metrics['sub'] = df_metrics['sub'].map(sub_map)

    df_metrics        = df_metrics.sort_values("feature").set_index("sub")
    df_metrics        = df_metrics.sort_values(input_metrics)
    metrics.append(df_metrics)

metrics = pd.concat(metrics).reset_index().sort_values(['sub_key', 'sub'])
metrics

```

```
[23]:
```

	sub	feature	sub_key	total_class	total_individuals \
0	0	hotel.0	hotel	1649	2111
1	1	hotel.1	hotel	2579	10307
3	0	house / apartment.0	house / apartment	3044	11727
2	1	house / apartment.1	house / apartment	743	691

		true_positive_rate	true_positive_rate_C
0	0.205584	0.020963	
1	0.302903	0.013879	
3	0.307458	0.013011	
2	0.162758	0.030811	

```
[24]: #metrics.reset_index().sort_values('sub_key')
metrics['color'] = [5,5, 2,2]
metrics['sub'] = ['hotel<br>False', 'hotel<br>True', 'house/apart<br>False', 'house/apart<br>True']
metrics = metrics.set_index('sub')
metrics
```

```
[24]:
```

	feature	sub_key	total_class \
sub			
hotel False	hotel.0	hotel	1649
hotel True	hotel.1	hotel	2579
house/apart False	house / apartment.0	house / apartment	3044
house/apart True	house / apartment.1	house / apartment	743

	total_individuals	true_positive_rate \
sub		
hotel False	2111	0.205584
hotel True	10307	0.302903
house/apart False	11727	0.307458
house/apart True	691	0.162758

	true_positive_rate_C	color
sub		
hotel False	0.020963	5
hotel True	0.013879	5
house/apart False	0.013011	2
house/apart True	0.030811	2

```
[25]: fig = plot_fairness_mistreatment(metrics, input_metrics, title="")

fig.update_layout(shapes=[
    dict(
        type= 'line',
        line=dict(
            width=1,
```



```

        dash="dot",
    ),
    xref='paper', x0= 0, x1= 1,
    yref='y',      y0= 1.5, y1= 1.5
)
])

# Set custom x-axis labels
fig.update_xaxes(nticks=4)

fig.update_layout(yaxis={'categoryorder':'array', 'categoryarray':['house/
    ↪apart<br>False', 'house/apart<br>True', 'hotel<br>False', 'hotel<br>True']})
fig.update_layout(xaxis_showgrid=False, yaxis_showgrid=False,
    ↪yaxis_title="Business")
fig.update_layout(font={'family': 'Courier New, monospace', 'size': 22},
    ↪height=550)

fig.show()
#fig.write_image("output/tmp/metrics/paper/recsys_fairnes_{}.eps".
    ↪format("business"))
#fig.show(renderer="svg")

```

Feature - Business Affinity

```

[26]: input_models_eval =
    ↪'InteractionTraining___mars_gym_model_b___logit_multipli_9dd8714dfd_96b9d686ee'
list_input_features    = ['childcare', 'family friendly']
input_metrics          = 'true_positive_rate'

```

```

[27]: metrics = []
for input_features in list_input_features:
    df_all_metrics    = load_fairness_metrics().loc[input_models_eval]
    df_instances      = load_fairness_df().loc[input_models_eval]
    df_all_metric_filter = df_all_metrics[df_all_metrics.sub_key.
    ↪isin([input_features])]

    columns           = list(np.unique(['sub_key', 'sub', 'feature',
    ↪'total_class', 'total_individuals'] + [input_metrics]))
    if input_metrics+"_C" in df_all_metrics.columns:
        columns.append(input_metrics+"_C")

    df_metrics        = filter_df(df_all_metrics, input_models_eval, columns,
    ↪'sub')

    df_metrics         = df_metrics[df_metrics.sub_key.isin([input_features])]
    if sub_map != None:

```

```

df_metrics['sub'] = df_metrics['sub'].map(sub_map)

df_metrics      = df_metrics.sort_values("feature").set_index("sub")
df_metrics      = df_metrics.sort_values(input_metrics)
metrics.append(df_metrics)

metrics = pd.concat(metrics)
metrics

```

```

[27]:
           feature      sub_key  total_class  total_individuals \
sub
1  childcare.1      childcare      1007      2490
0  childcare.0      childcare      3258      9928
0  family friendly.0  family friendly  2113      3313
1  family friendly.1  family friendly  2384      9105

           true_positive_rate  true_positive_rate_C
sub
1      0.237465      0.024281
0      0.287760      0.013512
0      0.200593      0.017856
1      0.311986      0.014863

```

```

[28]: metrics.reset_index()
metrics['sub'] = ['childcare<br>True', 'childcare<br>False', 'family_
↳friendly<br>True', 'family friendly<br>False']
metrics['color'] = [2,2, 5,5]
metrics = metrics.set_index('sub')
metrics

```

```

[28]:
           feature      sub_key  total_class \
sub
childcare<br>True      childcare.1      childcare      1007
childcare<br>False      childcare.0      childcare      3258
family friendly<br>True  family friendly.0  family friendly  2113
family friendly<br>False  family friendly.1  family friendly  2384

           total_individuals  true_positive_rate \
sub
childcare<br>True      2490      0.237465
childcare<br>False      9928      0.287760
family friendly<br>True  3313      0.200593
family friendly<br>False  9105      0.311986

           true_positive_rate_C  color
sub
childcare<br>True      0.024281      2

```

childcare False	0.013512	2
family friendly True	0.017856	5
family friendly False	0.014863	5

```
[29]: fig = plot_fairness_mistreatment(metrics, input_metrics, title="")

fig.update_layout(shapes=[
    dict(
        type= 'line',
        line=dict(
            width=1,
            dash="dot",
        ),
        xref='paper', x0= 0, x1= 1,
        yref='y',      y0= 1.5, y1= 1.5
    )
])

# Set custom x-axis labels
fig.update_xaxes(nticks=4)

fig.update_layout(yaxis={'categoryorder':'array', 'categoryarray':
    ↳['childcare<br>False', 'childcare<br>True', 'family friendly<br>False',
    ↳'family friendly<br>True']})
fig.update_layout(xaxis_showgrid=False, yaxis_showgrid=False,
    ↳yaxis_title="Business")
fig.update_layout(font={'family': 'Courier New, monospace', 'size': 22},
    ↳height=550)

fig.show()
#fig.write_image("output/tmp/metrics/paper/recsys_fairnes_{}.eps".
    ↳format("businesscare"))
#fig.show(renderer="svg")
```

```
[ ]:
```

Feature - City

```
[30]: input_models_eval =
    ↳'InteractionTraining___mars_gym_model_b___logit_multipli_9dd8714dfd_96b9d686ee'
input_features      = 'city_idx'
input_metrics       = 'true_positive_rate'
```

```
[31]: sub_map = {'0': "Boston",
                '1': "Como",
                '2': "Chicago",
                '3': "Lausanne",
                '4': "Dublin",
                '5': "Rio de Janeiro",
                '6': "Barcelona",
                '7': "Copenhagen",
                '8': "Vienna",
                '9': "New York",
                '10': "Hong Kong",
                '11': "Vancouver"}

fig, df_metrics, df_total = plot_fairness_metrics(input_models_eval,
↪input_features, input_metrics, sub_map = sub_map)
```

```
[32]: fig.show()
#fig.write_image("output/tmp/metrics/recsys_fairnes_{}.eps".
↪format(input_features))
```

```
[33]: df_metrics.head()
```

```
[33]:
```

	feature	sub_key	total_class	total_individuals	\
sub					
Como	city_idx.1	city_idx	34	30	
Dublin	city_idx.4	city_idx	265	1035	
Copenhagen	city_idx.7	city_idx	127	572	
Vienna	city_idx.8	city_idx	379	1282	
Barcelona	city_idx.6	city_idx	640	1652	
		true_positive_rate	true_positive_rate_C		
sub					
Como	0.184211		0.164835		
Dublin	0.252908		0.041650		
Copenhagen	0.289528		0.057278		
Vienna	0.291105		0.037876		
Barcelona	0.291691		0.029938		

```
[ ]:
```

Feature Device

```
[34]: input_models_eval =
↪'InteractionTraining___mars_gym_model_b___logit_multipli_9dd8714dfd_96b9d686ee'
df_instances = load_fairness_df().loc[input_models_eval]

input_features = 'device_idx'
```

```
input_metrics      = 'true_positive_rate'
input_items_top    = True
input_items        = []
```

[]:

```
[35]: fig = plot_fairness_treatment(df_instances, input_features, input_items,
    ↪top=input_items_top, title="", legend={"0": "Desktop", "1": "Mobile", "2":
    ↪"Tablet"})
fig.update_layout(font={'family': 'Courier New, monospace', 'size': 21},
    ↪height=600)

# Set custom x-axis labels
fig.update_yaxes(nticks=5)
fig.update_layout(xaxis_title="Hotel")

fig.show()
#fig.show(renderer="svg")
#fig.write_image("output/tmp/metrics/paper/recsys_fairnes_impact_{}.eps".
    ↪format(input_features))
```

/home/marlesson/anaconda3/lib/python3.7/site-
packages/numpy/core/_methods.py:234: RuntimeWarning:

Degrees of freedom <= 0 for slice

/home/marlesson/anaconda3/lib/python3.7/site-
packages/numpy/core/_methods.py:226: RuntimeWarning:

invalid value encountered in double_scalars

[]: